

# Synchronization and Bounded Buffer

# Semaphore

- Synchronization tool that provides more sophisticated ways (than Mutex locks) for process to synchronize their activities.
- Semaphore **S** – integer variable
- Can only be accessed via two indivisible (atomic) operations
  - **wait()** and **signal()**
    - Originally called **P()** and **V()**
- Definition of the **wait()** operation

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```
- Definition of the **signal()** operation

```
signal(S) {  
    S++;  
}
```

# Semaphore Usage

- **Counting semaphore** – integer value can range over an unrestricted domain
- **Binary semaphore** – integer value can range only between 0 and 1
  - Same as a **mutex lock**
- Can solve various synchronization problems
- Consider  $P_1$  and  $P_2$  that require  $S_1$  to happen before  $S_2$

Create a semaphore “**synch**” initialized to 0

**P1:**

$S_1;$

**signal(synch) ;**

**P2:**

**wait(synch);**

$S_2;$

- Can implement a counting semaphore  $S$  as a binary semaphore

# Semaphore Implementation

- Must guarantee that no two processes can execute the `wait()` and `signal()` on the same semaphore at the same time
- Thus, the implementation becomes the critical section problem where the `wait` and `signal` code are placed in the critical section
  - Could now have **busy waiting** in critical section implementation
    - But implementation code is short
    - Little busy waiting if critical section rarely occupied
- Note that applications may spend lots of time in critical sections and therefore this is not a good solution

## Semaphore Implementation with no Busy waiting

- With each semaphore there is an associated waiting queue
- Each entry in a waiting queue has two data items:
  - value (of type integer)
  - pointer to next record in the list
- Two operations:
  - **block** – place the process invoking the operation on the appropriate waiting queue
  - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue
- ```
typedef struct{  
    int value;  
    struct process *list;  
} semaphore;
```

## Implementation with no Busy waiting (Cont.)

```
wait(semaphore *S) {
    S->value--;
    if (S->value < 0) {
        add this process to S->list;
        block();
    }
}

signal(semaphore *S) {
    S->value++;
    if (S->value <= 0) {
        remove a process P from S->list;
        wakeup(P);
    }
}
```

# Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let  $s$  and  $q$  be two semaphores initialized to 1

| $P_0$                   | $P_1$                   |
|-------------------------|-------------------------|
| <code>wait(S);</code>   | <code>wait(Q);</code>   |
| <code>wait(Q);</code>   | <code>wait(S);</code>   |
| <code>...</code>        | <code>...</code>        |
| <code>signal(S);</code> | <code>signal(Q);</code> |
| <code>signal(Q);</code> | <code>signal(S);</code> |

- **Starvation – indefinite blocking**
  - A process may never be removed from the semaphore queue in which it is suspended
- **Priority Inversion** – Scheduling problem when lower-priority process holds a lock needed by higher-priority process
  - Solved via **priority-inheritance protocol**

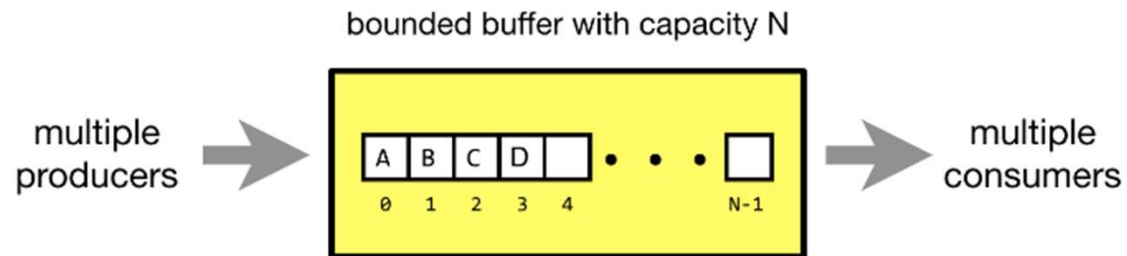
# Classical Problems of Synchronization

- Classical problems used to test newly-proposed synchronization schemes
  - Bounded-Buffer Problem
  - Readers and Writers Problem
  - Dining-Philosophers Problem



# Bounded-Buffer Problem

- $n$  buffers, each can hold one item
- Semaphore `mutex` initialized to the value 1
- Semaphore `full` initialized to the value 0
- Semaphore `empty` initialized to the value  $n$



# Bounded Buffer Problem


- A bounded buffer with capacity  $N$  has can store  $N$  data items. The places used to store the data items inside the bounded buffer are called **slots**.

Without proper synchronization the following **errors** may occur.

- The producers doesn't block when the buffer is full.
- A Consumer consumes an empty slot in the buffer.
- A consumer attempts to consume a slot that is only half-filled by a producer.
- Two producers writes into the same slot.
- Two consumers reads the same slot.

# Bounded Buffer

```
typedef struct {  
    int value[BUFFER_SIZE];  
    int next_in, next_out;  
} buffer_t;
```



Here,

- value used to store integer values in the buffer
- Index next\_in is used to keep track of where to write the next data item to the buffer.
- Index next\_out is used to keep track of from where to read the next data item from the buffer.

In the example, three data items B, C and D are currently in the buffer. On the next write data will be written to index next\_in = 4. On the next read data will be read from index next\_out = 0.

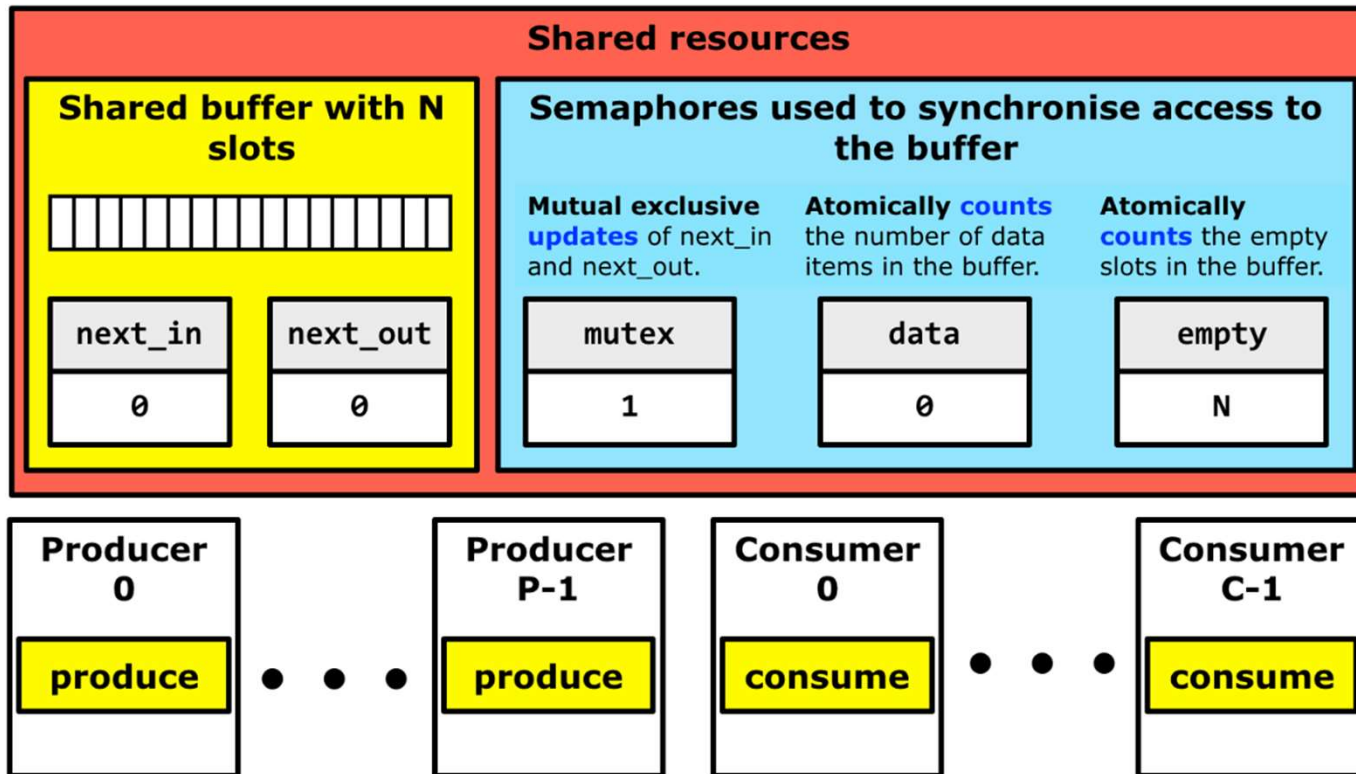
## Bounded Buffer - Critical Section and mutual Exclusion

- All updates to the buffer state must be done in a critical section. More specifically, mutual exclusion must be enforced between the following critical sections:
  - A producer writing to a buffer slot and updating next\_in.
  - A consumer reading from a buffer slot and updating next\_out.
  - A binary semaphore can be used to protect access to the critical sections.

# Synchronize producers and consumers

- Use one **semaphore** named **empty** to count the empty slots in the buffer. **Initialise** this semaphore to **N**.
  - A **producer** must **wait** on this semaphore before writing to the buffer.
  - A **consumer** will **signal** this semaphore after reading from the buffer.
- Use one **semaphore** named **data** to count the number of data items in the buffer. **Initialise** this semaphore to **0**.
  - A **consumer** must **wait** on this semaphore before reading from the buffer.
  - A **producer** will **signal** this semaphore after writing to the buffer.

# Bounded Buffer



# Bounded Buffer Problem (Cont.)

- The structure of the producer process

```
do {  
    ...  
    /* produce an item in next_produced */  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    /* add next produced to the buffer */  
    ...  
    signal(mutex);  
    signal(data);  
} while (true);
```

# Bounded Buffer Problem (Cont.)

- The structure of the consumer process

```
Do {  
    wait(data);  
    wait(mutex);  
    ...  
    /* remove an item from buffer to next_consumed */  
    ...  
    signal(mutex);  
    signal(empty);  
    ...  
    /* consume the item in next consumed */  
    ...  
} while (true);
```



# Readers-Writers Problem

- A data set is shared among a number of concurrent processes
  - Readers – only read the data set; they do **not** perform any updates
  - Writers – can both read and write
- Problem – allow multiple readers to read at the same time
  - Only one single writer can access the shared data at the same time
- Several variations of how readers and writers are considered – all involve some form of priorities
- Shared Data
  - Data set
  - Semaphore `rw_mutex` initialized to 1
  - Semaphore `mutex` initialized to 1
  - Integer `read_count` initialized to 0

